The deadline for this exercise is on Friday 25.06.2021 at 23:59

Structure from Motion

- You will need to install the following libraries for this exercise

  `'pip install tqdm'`

  When you are finished with this exercise,

- Compress the complete directory into a ZIP and upload it to StudOn. Make sure folder *results/ex3* exists in the project root. On Unix (and powershell) systems, the following command generates the submission archive

  ```
  cd ComputerVision2021 (or CV2021);
  # as a single line:
  zip ../submission_ex3.zip ./data/ex3/*.jpg ./bin/run_ex3.py
  ./ex3/*py test/test_ex3.py ./results/ex3/*png
  ```

- For groups, one member uploads the submissions and adds their partner to the exercise. Don't forget to add all members of the team in *ex3/__init__.py*

- We are not grading based on vectorization in this exercise. However, your code should not take more than *2 minutes* to run for your own convenience.

1 Introduction

In this exercise, we will compute the relative transformation of two cameras, i.e. rotation and translation, and the 3D position of a set of feature matches. Consider Fig. 1 that shows a *memorial shield* captured from two different angles. Given a point x in image 1 and the corresponding match x' in image 2, the following constraint can be derived from epipolar geometry (see lecture):

$$\mathbf{x}'^{T}\mathbf{F}\mathbf{x} = 0, \tag{1}$$

The 3×3 matrix $\mathbf{F}$ is the fundamental matrix, mapping points in one image to epipolar lines in the other image.

Computer Vision Exercise   Angelos Nikolaou

Exercise 3   Prathmesh Madhu

Structure from Motion   10.06.2021

LGDV

COMPUTER GRAPHICS • ERLANGEN

LME

Figure 1: A *memorial shield* captured from two different viewpoints at the Germanic National Museum.

2 Computation of F

Similar to the homography computation of the last exercise sheet, the fundamental matrix can be computed from 8 matching feature points. Look out for an additional step for fundamental matrix as compared to homography. Also like previous times, this procedure is wrapped in a RANSAC algorithm (this time the function is provided to you) to account for incorrect matches.

2.1 Distance from the Epipolar Line

Implement the function `inliers_epipolar_constraint` in `ex3/functions.py` that checks, if a given match satisfies the epipolar constraint. This is the case, if the distance of x' to the epipolar line of x is smaller than a threshold t.

- Compute the epipolar line $l = (l_x, l_y, l_w) = Fx$

- Normalize l to get the scale independent line $l' = l/\sqrt{l_x^2 + l_y^2}$

- Compute the point line distance $d = x'^T l'$

- A match is an inlier if $|d| < t_e$

2.2 8-Point Algorithm

Given 8 correct feature matches, the fundamental matrix can be computed with the 8-point algorithm. Implement the function `compute_fundamental_matrix` in `ex3/functions.py` with the following steps:

1. Construct the matrix A. (Note: 1 row per match $\rightarrow$ 8 rows in total)

$$A = \begin{bmatrix} p_{x1}q_{x1} & p_{x1}q_{y1} & p_{x1} & p_{y1}q_{x1} & p_{y1}q_{y1} & p_{y1} & q_{x1} & q_{y1} & 1 \\ p_{x2}q_{x2} & p_{x2}q_{y2} & p_{x2} & p_{y2}q_{x2} & p_{y2}q_{y2} & p_{y2} & q_{x2} & q_{y2} & 1 \\ \dots & & & & & & & & \\ \dots & & & & & & & & \\ p_{x8}q_{x8} & p_{x8}q_{y8} & p_{x8} & p_{y8}q_{x8} & p_{y8}q_{y8} & p_{y8} & q_{x8} & q_{y8} & 1 \end{bmatrix}$$

2. Solve $Af = 0$ and extract F

$$f = \begin{bmatrix} f_0 & f_1 & f_2 & f_3 & f_4 & f_5 & f_6 & f_7 & f_8 \end{bmatrix}^T$$

$$F = \begin{bmatrix} f_0 & f_3 & f_6 \\ f_1 & f_4 & f_7 \\ f_2 & f_5 & f_8 \end{bmatrix}$$

3. Make sure that $Rank(F) = 2$.
   - Compute the SVD $F = U\Sigma V^T$ with

$$\Sigma = \begin{bmatrix} \sigma_1 & 0 & 0 \\ 0 & \sigma_2 & 0 \\ 0 & 0 & \sigma_3 \end{bmatrix}$$

   - Set $\sigma_3 = 0$
   - Recompute F with the updated Σ.
   - Normalize F.

There's a test-case in *test/test_ex3.py* to validate your 8-point implementation. This test should succeed before continuing with the next tasks.


Figure 2: The *geometric consistent* matches. (left) All matches **before** passing the *ratio test* **and** satisfying the epipolar constraint; (right) All matches **after** passing the *ratio test* **and** satisfying the epipolar constraint

3 Triangulation

Triangulation describes the method of finding the position of a point in 3D space given its image in two views and the camera parameters of both views (see Figure 3).


Figure 3: A point $\boldsymbol{x}$ is seen by two cameras. The rays starting from each camera center passing through the imaged points $\boldsymbol{p}$ and $\boldsymbol{q}$ intersect at the world point $\boldsymbol{x}$.

Given the pinhole camera model with the projection matrices $M_i = K_i V_i$ the projection equations

$$\boldsymbol{p} = \boldsymbol{M_1 x}$$

$$\boldsymbol{q} = \boldsymbol{M_2 x},$$

can be transformed to the homogeneous linear system of equations $Ax = 0$ with

$$\boldsymbol{A} = \begin{bmatrix} p_x\boldsymbol{m}_1^3 - \boldsymbol{m}_1^1 \\ p_y\boldsymbol{m}_1^3 - \boldsymbol{m}_1^2 \\ q_x\boldsymbol{m}_2^3 - \boldsymbol{m}_2^1 \\ q_y\boldsymbol{m}_2^3 - \boldsymbol{m}_2^2 \end{bmatrix},$$

where $\boldsymbol{m}_i^j$ is the j-th row-vector of M_i.

Implement the function `triangulate` in `ex3/functions.py` with the description from above. Check if the test case, which is executed in `test/test_ex3.py`, succeeds before you continue with the exercise.

4 Relative Transformation

It is shown that the camera positions are mathematically related via Essential matrix: $E = K'^T F K$; where K and K' are the intrinsic camera parameter matrix of camera 1 and 2 respectively. In our case, because we are using one camera and move it to generate a second image, $K = K'$. Therefore we have the camera position 1 set at point (0; 0; 0), the relative position of camera 2 can be determined (up to scaling).

4.1 Computing Essential Matrix

So then to compute the relative transformation (rotation and translation) between the input images, first, the essential matrix E is computed from the fundamental matrix F and the camera intrinsics K:

$$E = K^T F K$$

Implement the function `compute_essential_matrix` in `ex3/functions.py`. Also, there's a test to check if the essential matrix is correct in `test/test_ex3.py`.

4.2 Decomposing Rotation and Translation Matrices

The essential matrix is then decomposed into two rotation matrices and translation vectors, giving a total number of 4 possible transformations per essential matrix. For each relative transformation, the feature matches are triangulated (Task 3) and projected to each image. Only one solution over all pose configurations is geometrically valid, which means that all 3D points lie in front of both cameras. The other configurations have at least one point that is behind one or both cameras.

Implement the function `decompose` in `ex3/functions.py` with the following steps:

1. Compute the SVD $E = U \Sigma V^T$

2. The two possible rotations are

$$R_1 = U W V^T$$

$$R_2 = U W^T V^T$$

with

$$W = \begin{bmatrix} 0 & -1 & 0 \\ 1 & 0 & 0 \\ 0 & 0 & 1 \end{bmatrix}$$

3. The two possible translations are

$$t_1 = s$$

$$t_2 = -s$$

with

$$s = u_3/|u_3|$$

and u_3 being the third column vector of U.

There's a test to check the decompose function generates correct rotation and translation matrices in `test/test_ex3.py`. The View 2 matrix after relative Transformation that uses `decompose` internally will not be exact for everyone but for reference should look like this:

```
View2:
[[ 0.93126  -0.15450   0.32997  -0.99363]
 [ 0.17217   0.98475  -0.02482  -0.11094]
 [-0.32111   0.07993   0.94366   0.01979]]
```

Point cloud Visualization

We wanted to visualize the point clouds that we compute as seen in Fig. 3. All the functions have been implemented for you. This part is ofcourse not graded but it very important to understand and visualize the concepts. The 3D visualization should look similar as below:

Figure 4: The two view reconstruction and final output after task 4. We have computed the relative transformation between both views and a set of 3D points that are seen by both cameras. The blue and the red tips represent the camera positions and the points in gray represents a 3d pointcloud, i.e. the pixels of images from Fig. 1. These 3 images are representing the same figure from 3 different angles. Open3d lets you rotate this in your system in anyway you want for better understanding.